

Per-scenario detection over 15 runs (mean $\pm$ 95% CI). Red = attack (recall, higher better); blue = benign (FPR, lower better)

